

ContractOS

AI-Powered Contract & Payroll Intelligence
for Oracle Game Studio

Daria Bystrova & Julian Granados · SilverTrust · Pitch to Review Board



01 Client & Problem

02 The Solution

03 Compliance

04 Monitoring

05 The Ask

Our Recommendation

OVERVIEW

Oracle's burn crisis is a cost visibility problem — ContractOS solves it with five AI agents that read every contract, catch every invoice error, and show exactly where the money is going.

Three supporting arguments

01

The problem is real

9 pain points discovered blind. Revenue side cannot be fixed. Only lever is cost.

02

The solution is defensible

5 AI agents, one per capability. Human-in-the-loop technically enforced. EU AI Act minimal/limited risk.

03

The system is observable

LangSmith live — 6 runs complete, all 5 agents traced, PII redacted, full audit trail for regulators.

Oracle has no visibility over where the money goes

Discovered blind — 10 indirect questions, 0 direct pain-point asks

01 CLIENT & PROBLEM

900 people · 5 jurisdictions · high burn · no contract register · no single cost view.

CEO / CTO: Eugen

CFO: Thibaud

Size: ~900 FTEs, part-time, contractors, interns

Countries: Russia · India · Pakistan · Bangladesh · Philippines

Revenue: B2C games, one-time purchase, no subscriptions

Lever: Revenue side is structurally fixed — only cost can move

Pain points uncovered

- No contract visibility — 900 contracts unread
- Wrong invoicing — amounts not matching contracts
- Lost contractors — duplicate or inactive payments
- Paying twice — contractor = internal role scope
- No cost visibility — burn untraceable by category
- Multi-country payroll bottleneck + intermediaries
- Compliance unknown — asked unprompted
- IP at risk — ambiguous contractor IP clauses

Every pain point was earned through indirect questioning

01 DISCOVERY EVIDENCE

No direct asks. These verbatim quotes from the transcript confirm the problem is real.

Eugen *"I cannot really tell you what kind of tools we're using... people just want a tool and then we give it to them."* → **Zero cost visibility**

Thibaud *"The stack is absolutely a mess. Wrong invoicing, pay cycles. There are a lot of problems."* → **Core pain confirmed**

Thibaud *"We kind of lost track of all the subcontractors we have and we're working with."* → **Duplicate payments likely**

Both *"Are we actually compliant by doing what we are doing?" — asked unprompted.* → **Compliance anxiety real**

ContractOS — five agents, one per capability, zero automated decisions

02 THE SOLUTION

One interface. Five AI agents orchestrated by LangGraph. Every output goes to a human before any action.

1

Contract Intelligence

Agent 1

- LLM ingests contract documents.
- Structured extraction prompt returns JSON: name, role, rate, currency, jurisdiction, notice period, IP clause status, GDPR clause status, working arrangement description.
- Output stored in Contract Register DB.

2

Invoice Validation

Agent 2

- LLM reads incoming invoice text and compares it against extracted contract terms.
- Flags discrepancies in amount, currency, scope, and billing period.
- Output: validated / flagged + specific mismatch detail.

3

Overlap Detection

Agent 3

- LLM compares contractor working arrangement descriptions against internal job role descriptions using semantic similarity.
- Flags pairs where scope overlap exceeds threshold.
- Output: overlap alert with contractor name, internal role, estimated duplicate cost.

4

Compliance Flagging

Agent 4

- LLM checks extracted contract data against jurisdiction-specific rule sets: GDPR clause presence, IP assignment clarity, misclassification indicators, notice period minimums. Output: Red / Amber / Green status + specific flags per contract.

5

Cost Categorisation

Agent 5

- LLM normalises ambiguous spend descriptions and categorises all payments by country, contract type, team, and intermediary.
- Plain software aggregates and renders the dashboard.
- Output: categorised cost register + burn trend.

⚠ Human approves every flag before any payment is sent or any contractor is reclassified — this is technically enforced, not just policy.

AI only where language understanding is required — everything else is plain software

02 AI vs SOFTWARE

This boundary is what keeps the system auditable, maintainable, and cost-efficient.

Needs LLM reasoning

- ◆ Extracting data from unstructured multilingual contract templates
- ◆ Semantic overlap: contractor scope vs internal job role descriptions
- ◆ Misclassification risk from natural language working arrangements
- ◆ Missing clause detection using jurisdiction-specific legal knowledge
- ◆ Invoice anomaly detection across inconsistent multilingual formats

Plain software sufficient




- Structured database storage of extracted fields
- Cost dashboard charts and aggregation
- Payment routing via Wise API or local rails
- Role-based access control
- Audit trail logging and notifications

Sample Dashboard by Plain Software Integration (Demo Only!)


02 AI vs SOFTWARE

ContractOS — Agent 5: Cost Categorisation Output

Oracle Game Studio - Monthly burn: €96,900 | 20 payment records processed



Intermediary cost: €45,500 (47% of total burn)



Agent 5 flags

⚠ Overlap detected

Jana Novak — €8,500/month
95/100 overlap with internal engineer

⚠ Intermediary opacity

€45,500/month (47%)
Via agencies — no individual contracts

✓ Categorisation complete

20 records processed
confidence avg: 92%

→ Recommended action

Review 4 intermediary accounts
Terminate Jana duplicate — save €8,500/mo

⚠ Intermediary opacity risk — who are the underlying contractors?

ContractOS - Agent 5 (Cost Categorisation Agent) - Output: categorised cost register + burn trend - SilverTrust Project 4 - Daria Bystrova & Julian Granados

Total monthly burn: €96,900

ContractOS is minimal/limited risk — justified, not asserted

03 COMPLIANCE — EU AI ACT

No agent executes a decision. Every output goes to a named human before action. This is what keeps the classification defensible.

MINIMAL RISK

Contract extraction · Cost dashboard · Payment tracking

Surfaces information only. No decisions affecting individuals.

LIMITED RISK

Overlap detection · Misclassification flagging

Touches Annex III pt.4 territory — human legal review required before any reclassification.

Roles & Obligations

SilverTrust = PROVIDER

- Maintain technical documentation
- Implement quality management + logging
- Ensure human oversight is technically enforced — not just policy

Oracle Game Studio = DEPLOYER

- Maintain human review for all flags before action
- Inform workers AI analyses their contracts (transparency obligation)
- Conduct own DPIA before deployment

Every data use has a named lawful basis — three LIAs completed

03 GDPR & OTHER LAW

No personal data leaves EU infrastructure. LLM processing runs under documented legitimate interest.

Personal data use	Lawful basis	Article	Constraint / Action
Employment contracts	Contract performance	6(1)(b)	No consent needed
Payroll records & invoices	Legal obligation	6(1)(c)	Tax + labour law retention
LLM contract processing	Legitimate interest	6(1)(f)	LIA completed · EU infra only · DPA signed
LangSmith traces	Legitimate interest	6(1)(f)	PII redacted · 90-day retention max

Other Applicable Law

NL DBA Act (2025): Agent 4 flags misclassification — human + legal review before reclassification

German BetrVG: Works council consulted before go-live — SilverTrust provides technical documentation

India DPDP / Philippines DPA: Data stays EU-side — local counsel for processing obligations

Russia: Out of scope v1 — sanctions + data localisation require specialist legal advice

The black box has windows — every AI decision is logged and auditable

04 MONITORING

If a regulator asks what happened with a specific contractor's data, we show the complete chain in under five minutes.

Extraction confidence Alert < 0.80 on rate, jurisdiction, IP

Failed extractions Zero tolerance — immediate alert

Human override rate Alert > 20% in 7-day window

Compliance flag trend Alert if Red > 10% of new contracts

Latency per agent Alert p95 > 30 seconds

PII redaction audit Weekly — zero tolerance for PII in traces

✓ **Live — contractos-oracle-prod**

6 runs complete · All 5 agents traced · EU endpoint: eu.api.smith.langchain.com · PII redacted

"Here is where you can see everything the AI did, catch a mistake before it costs you money, and prove to an auditor exactly what happened — in under five minutes."

05 The Ask

05 THE ASK

- ✓ Problem discovered blind — 9 pain points confirmed by verbatim transcript evidence
- ✓ Solution addresses all 9 pain points with 5 dedicated AI agents — each maps to exactly one capability
- ✓ EU AI Act: Minimal / Limited risk — justified, not asserted — human-in-the-loop technically enforced
- ✓ GDPR: 3 lawful bases documented, 3 LIAs completed, subject rights handled, data stays EU-side
- ✓ LangSmith: 6 runs live, all 5 agents traced, PII redacted, complete audit trail
- ✓ ContractOS gets Oracle back to black numbers — cost reduction measurable from day one

Do you approve ContractOS for Oracle Game Studio — or what needs to change?

ContractOS v1.1

Revised Proposal — Based on Client Feedback

Oracle approved ContractOS v1.0 — seven change requests incorporated below.

CR-01 EOR coverage

CR-02 Social registration

CR-03 Contract generation

CR-04 PO matching

CR-05 Accounting rules

CR-06 Intermediary costs

CR-07 Multilingual QA

Oracle approved — with seven change requests

06 REVISED PROPOSAL

ContractOS v1.0 architecture, compliance position, and LangSmith monitoring approved without objection. Seven enhancements requested during Q&A.

Approved unchanged ✓

- Five-agent architecture (Agents 1–5)
- Human-in-the-loop — technically enforced
- EU AI Act: Minimal / Limited risk
- GDPR: 3 lawful bases, 3 LIAs completed
- LangSmith observability live
- Auto-signing rejected — human signature always required

Change requests → v1.1

- CR-01 EOR coverage per country **Phase 1**
- CR-02 Social registration automation **Phase 2**
- CR-03 Contract generation workflow **Phase 2**
- CR-04 PO matching — Agent 2 **Phase 1**
- CR-05 Local accounting rules — Agent 2 **Phase 1**
- CR-06 Intermediary cost tracking **Phase 1**
- CR-07 Multilingual QA plan **Phase 1**

Phase 1 — immediate enhancements to existing agents

06 REVISED PROPOSAL

Four change requests implemented as enhancements to Agent 2, 4, and 5. No new agents required. Ready for v1.1.

Agent 4

EOR Coverage

CR-01

Flags contracts in countries where Oracle has no legal entity (Russia, India, PK, BD, PH). Output: EOR required flag + estimated EOR cost range + recommended action. Agent 5 tracks EOR fees as a named cost category.

Agent 2

PO Matching + Local Accounting Rules

CR-04 + CR-05

Agent 2 now compares invoice against contract AND Purchase Order (three-way match). Also checks jurisdiction-specific accounting thresholds — Germany, India, Philippines in v1.1. Flags if invoice triggers different accounting treatment.

Agent 2 + 5

Intermediary Cost Tracking

CR-06

Intermediary fees (staffing agencies, EOR providers, government fees) tracked as adjacent costs per contract. New cost subcategories: `staffing_agency` · `eor_provider` · `government_fee` · `legal_counsel`. Alert if intermediary costs exceed 30% of country spend.

Phase 2 — new capabilities after go-live stabilisation

06 REVISED PROPOSAL

Client explicitly requested full process automation up to the human signature step. Two Phase 2 capabilities scoped — compliance research required before commitment.

CR-02 Social Registration Automation

Trigger: New contract created in target country

- Pulls country-specific social security + insurance registration requirements
- Pre-fills registration forms with contract data
- Routes to HR for review and official submission
- Human submits to government bodies — agent prepares only

⚠ **Pre-condition: legal review per jurisdiction before automation**

CR-03 Contract Generation Workflow

Trigger: HR onboards a new worker

- Collects: name, role, country, contract type, rate, start date
- LLM generates draft contract using jurisdiction-specific template
- Draft produced in local language + English
- Lawyer reviews first template per country — approved templates reused

⚠ **Auto-signing explicitly out of scope — human signature always required**

Multilingual QA plan + revised scope summary

CR-07: Formal three-phase QA plan for Agent 1 across all operating languages. Revised scope table with phases and priorities.

CR-07 — Multilingual QA Plan

Phase	Action	Acceptance
Phase 0	5 test contracts per language, manual verification	≥90% field accuracy
Phase 1	Refine extraction prompts per language from errors	≥95% field accuracy
Phase 2	LangSmith monitors confidence scores in production	Alert if < 0.80

Languages: English · Russian · Hindi · Urdu · Filipino · Bengali · German

Revised scope summary

Capability	Phase	Priority
Five-agent core (v1.0)	Delivered	✓
EOR flagging (Agent 4)	v1.1 — now	●
Intermediary tracking (A2+A5)	v1.1 — now	●
PO matching (Agent 2)	v1.1 — now	●
Local accounting (Agent 2)	v1.1	●
Multilingual QA plan	v1.1	●
Social registration auto	Phase 2	●
Contract generation workflow	Phase 2	●
Auto-signing	Out of scope	✗

07 The Revised Ask

07 THE REVISED ASK

- ✓ ContractOS v1.0 approved — five agents, compliance position, LangSmith monitoring
- ✓ Seven change requests received and documented — all addressed in this revised proposal
- ✓ Phase 1 enhancements ready: EOR flagging · PO matching · intermediary cost tracking · accounting rules
- ✓ Phase 2 scoped: social registration automation + contract generation workflow (post go-live)
- ✓ Multilingual QA plan formalised: 3 phases, acceptance criteria per language, 7 languages covered
- ✓ Auto-signing out of scope — human signature always required — accepted by client

Do you approve ContractOS v1.1 for Oracle Game Studio?